

Across Africa, universities, startups, NGOs and public institutions are building AI tools for their own communities: clinic helplines, school assistants, customer support in local languages. Many lack a dedicated safety layer, because the standard options — a second guardrail model or a per-message moderation API — are priced beyond these organisations. The users left exposed write in Nigerian Pidgin, Hausa or Igbo, which imported filters serve worst anyway. Our project investigates whether the signals needed to identify harmful requests are already present inside a language model's own internal representations, so safety can be added at negligible cost and remain open to inspection. Guided by equity, accessibility, transparency and proportionality, we will report openly where our approach succeeds and where it falls short, including on under-resourced African languages where appropriate to our dataset. Responsible AI should not depend on an organisation's budget: effective safety should be within reach of every African builder.