

STAKEHOLDER ENGAGEMENT PLAN

Latent Probing for Toxicity Detection

Team Name: Team OKU | Team Leader: Thaabit Idris
Team Members: Ogwu Naomi Chukwufumnanya, Riyas Yusuf
Cohort: AI Saturdays Lagos - Cohort 10
Competition: Latent Probe Challenge - Toxicity Detection (Codabench)
Date: September 2026

1. Stakeholder Groups

Group 1 - Deployers (African AI builders)

Who they are: Small development teams, university labs, NGOs, and public institutions building AI assistants for clinics, schools, and public services across Africa. Why they matter: they are the direct adopters of our probe - without their buy-in, the tool never reaches an end user at all. What matters most to them: low integration cost, a low false-positive rate (a tool that blocks legitimate messages breaks their product), and transparency into why something was flagged, since most cannot afford to debug an opaque third-party system.

Group 2 - End Users (Nigerian Pidgin, Hausa, Igbo, and Yoruba speakers)

Who they are: The people typing messages into assistants built by Group 1 - patients using a clinic helpline, students using a school assistant, citizens using a public information service. Why they matter: they are the ultimate subject of every decision our probe makes, and our own Data Card documents that we do not yet know how the probe behaves on their language. What matters most to them: not having ordinary speech in their own language wrongly blocked as toxic, and trusting that a safety tool actually protects them rather than silently misjudging how they write.

We prioritised this pairing deliberately: Group 2 is precisely the population our Data Card names as unverified by the current English-only probe, so engaging them is the mechanism that would surface whether that documented gap can genuinely be closed.

2. Potential Benefits and Risks

Group 1 - Deployers

Benefits: a near-free safety layer they currently cannot afford, reusing computation their model already performs. Risks: if they deploy assuming broader language coverage than the probe actually has, they inherit a false sense of safety for their non-English-speaking users. Economic consideration: adopting our probe without also verifying it on their actual user base could create a worse outcome than having no probe and knowing it.

Group 2 - End Users

Benefits: once extended to their languages, a low-cost safety layer protecting them in interactions currently unscreened. Risks: right now, specifically, there is no benefit to this group, only exposure - a deployer using our probe on non-English input is applying an untested tool to their messages. Cultural consideration: even after language coverage is added, the toxic-class definition is rooted in Western safety norms (documented in our Data Card) and may not reflect what this group itself considers harmful or acceptable speech.

3. Engagement Method

Group 1 - Deployers: User Committees

We selected User Committees for the deployer group because their relationship to this project is ongoing and operational, not a single decision point. A deployer does not evaluate our probe once and walk away - they integrate it, watch it run against real user traffic for weeks or months, and encounter failure modes (a false positive blocking a legitimate clinic query, a false negative letting a harmful request through) that only surface in production. A one-time consultation would capture their expectations before deployment but miss the far more valuable feedback that only appears after deployment. A standing committee structure matches this reality: it creates a recurring touchpoint where deployers bring concrete examples from their own systems rather than hypothetical concerns.

The method also fits the size and nature of this group. Deployers are a small, identifiable, and motivated population - teams who have actively chosen to integrate a safety tool, meaning they already have both the technical context and the incentive to engage seriously. This is precisely the condition under which a committee structure works well: it does not require recruiting or persuading disengaged participants, only convening people who are already invested.

In practice, we would convene a small recurring group of three to five deployer teams, meeting on a regular cadence (proposed quarterly, adjustable based on how quickly integration issues surface). Each session would review a sample of real flagged and unflagged examples from their production use, discuss disagreements between the probe's output and the deployer's own judgment, and propose concrete changes - to documentation, to the abstain-band threshold described in our Impact Statement, or to how flagged decisions are explained to their own end users. Committee members would be the first group consulted before any change to the probe's training data or threshold is finalised, giving them real influence rather than a symbolic consultative role.

This method directly serves the values named in our Data Card: transparency, because committee members see actual flagged examples rather than aggregate statistics; proportionality, because their feedback is the most direct signal we have for whether the probe suppresses legitimate speech in their specific deployment context; and accountability, because a standing committee - unlike a one-time survey - gives deployers a channel to hold us to the limitations we have documented, rather than trusting a static document alone.

Group 2 - End Users: Distributed Dialogue

We selected Distributed Dialogue for the end-user group because this population cannot be reached the way the deployer group can. End users are geographically dispersed across multiple countries and language communities, do not share a single platform or point of contact, and have no existing relationship with our team - unlike deployers, who opt in by choosing to integrate our probe. A centralised method, such as a single workshop or committee, would only reach whichever community happened to be reachable by our team directly, most likely English-speaking participants in or near Lagos. That would reproduce the exact bias our Data Card already names: overrepresenting the easiest-to-reach population while leaving Pidgin, Hausa, Igbo, and Yoruba speakers - the groups our problem statement is actually about - underrepresented in the very process meant to correct their underrepresentation.

Distributed Dialogue is designed for exactly this situation: decentralised conversations conducted through multiple local channels rather than one central venue, allowing engagement to happen where each community already is rather than requiring them to come to us. This also respects a practical reality - genuine trust with language communities who have historically been served worst by imported technology is not built through a single outside-initiated session, but through sustained contact via intermediaries the community already trusts.

In practice, we would partner with existing community and language organisations already working with Pidgin, Hausa, Igbo, and Yoruba speakers, rather than recruiting participants directly ourselves. Each partner organisation would run conversations in their own community, in the relevant language, surfacing examples of what their community considers harmful or acceptable speech, and specifically what kinds of messages they would be uncomfortable seeing wrongly flagged or wrongly allowed through. These conversations would be conducted independently across communities rather than merged into a single averaged summary, mirroring the per-language reporting principle already established in our Data Card, so that, for example, Yoruba-specific input is not diluted by input from the other three languages.

This method is the only realistic path to closing the linguistic-fairness gap our Data Card names as currently unmet. It cannot be closed by translation alone, since translation without community input risks encoding an outsider's assumptions about what counts as harmful in each language. Distributed Dialogue is what would let this project move from documenting the gap honestly to actually closing it.

4. Influence on Project Design

Deployer feedback would most directly influence evaluation and deployment decisions - real production false-positive patterns would inform our abstain-band threshold and per-language accuracy reporting. End-user feedback, once Distributed Dialogue is underway, would most directly shape dataset and model design: it is the only route to closing the linguistic-fairness gap documented in our Data Card, since it would surface which harmful-request patterns in each language actually need representation, rather than our team guessing through translation alone. Both groups' input would refine outputs - deployers toward clearer flagging explanations, end users toward ensuring flags reflect locally-relevant definitions of harm rather than imported ones.